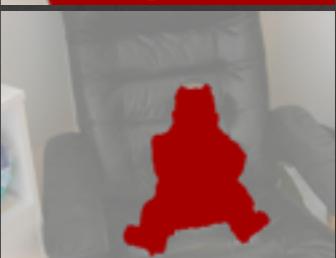


Geodesic Object Proposals

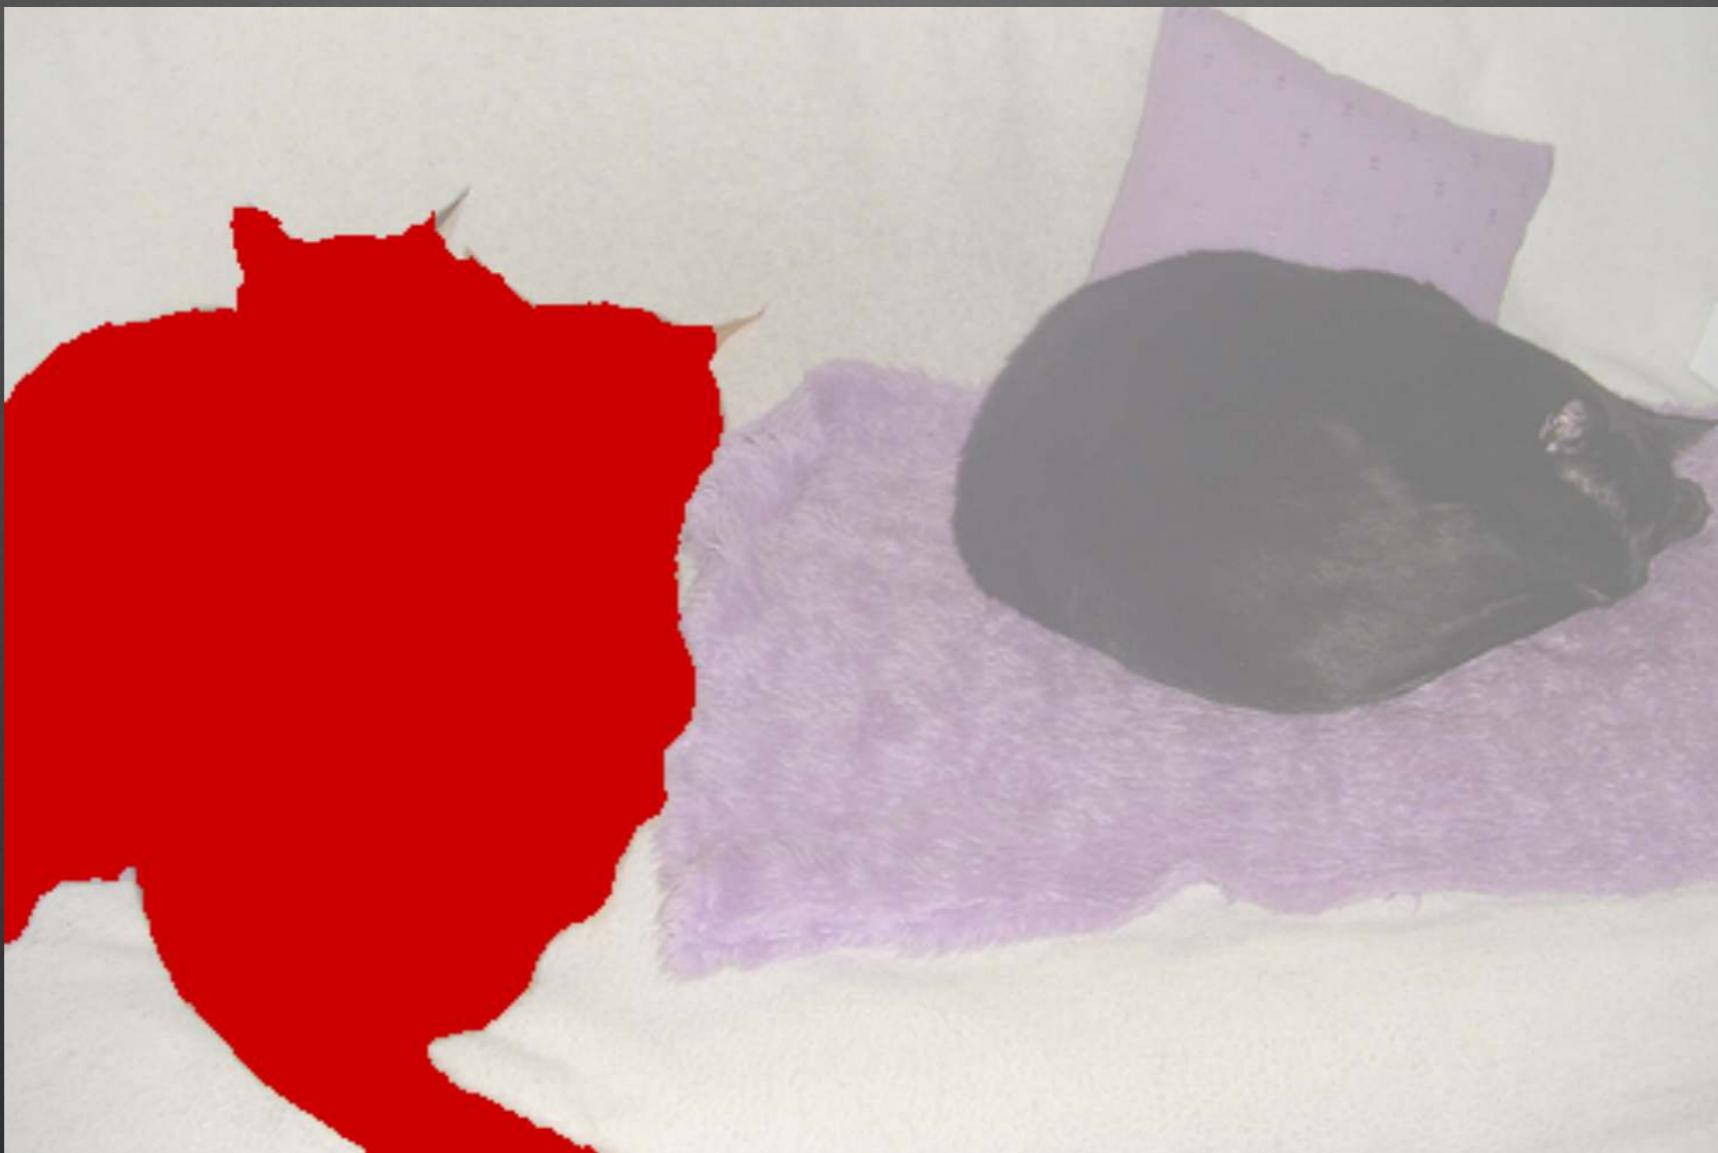
Philipp Krähenbühl
Stanford University

Vladlen Koltun
Adobe Research



Object Proposals

Find small set of proposals
that includes all objects in a scene



Object Proposals

Bounding Box
Proposals

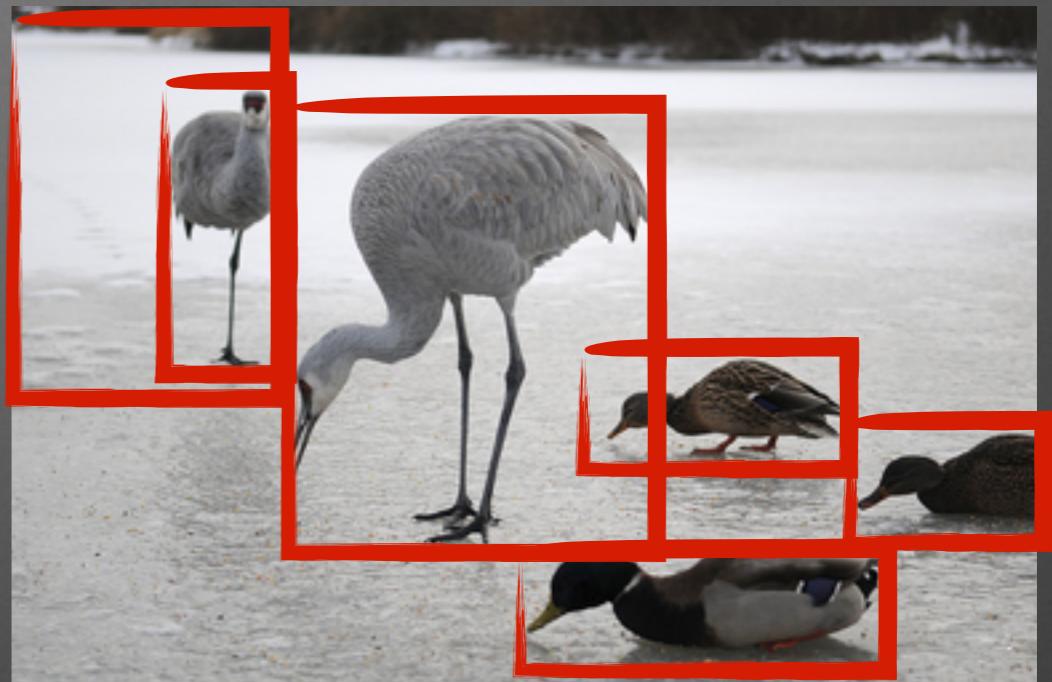


Segment
Proposals



Uses of object proposals

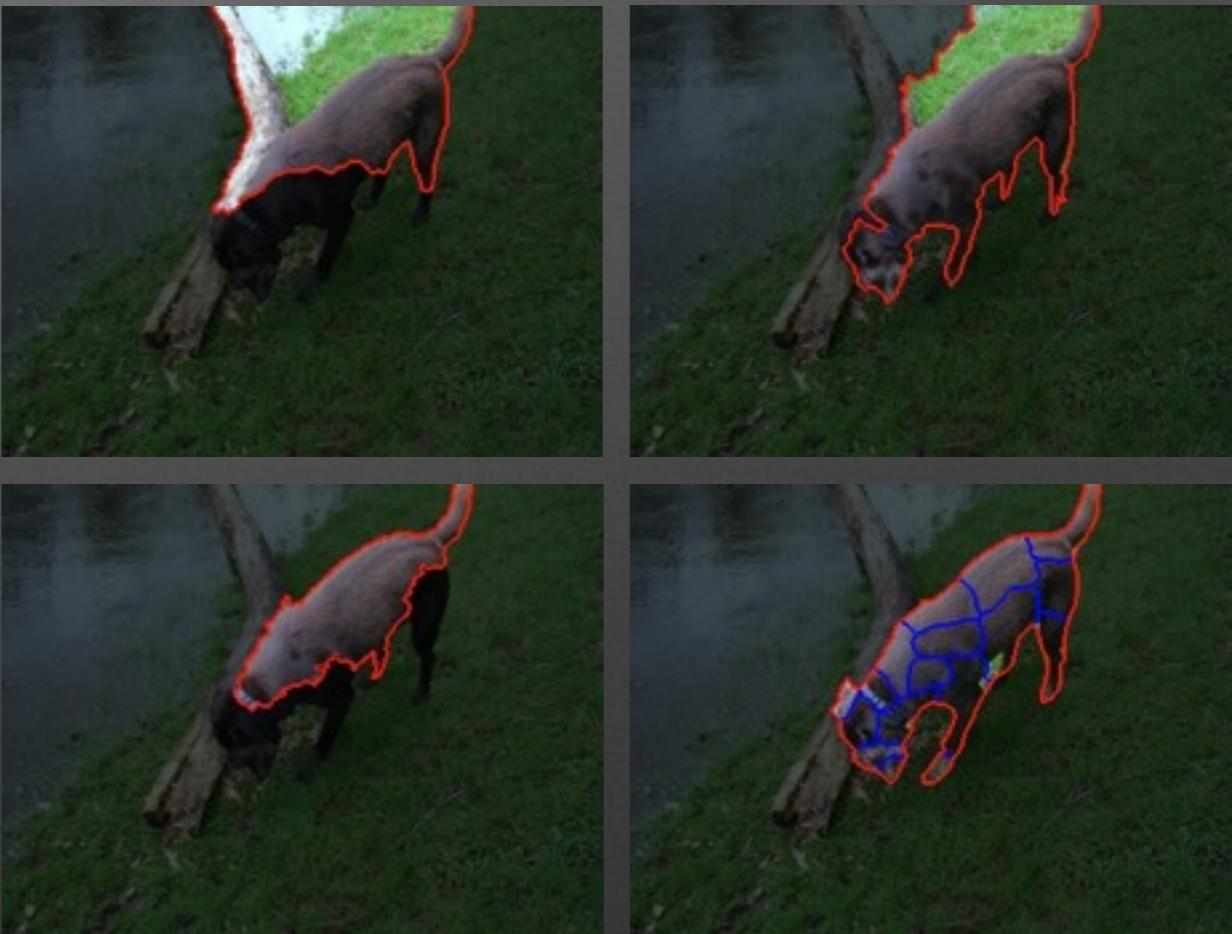
- Object detection
 - faster
 - more expensive features
 - used in almost all state-of-the art methods
- Multi-class image segmentation
 - used in state-of-the art
 - more expressive features



Improving Spatial Support for Objects via Multiple Segmentations

[Malisiewicz and Efros 2007]

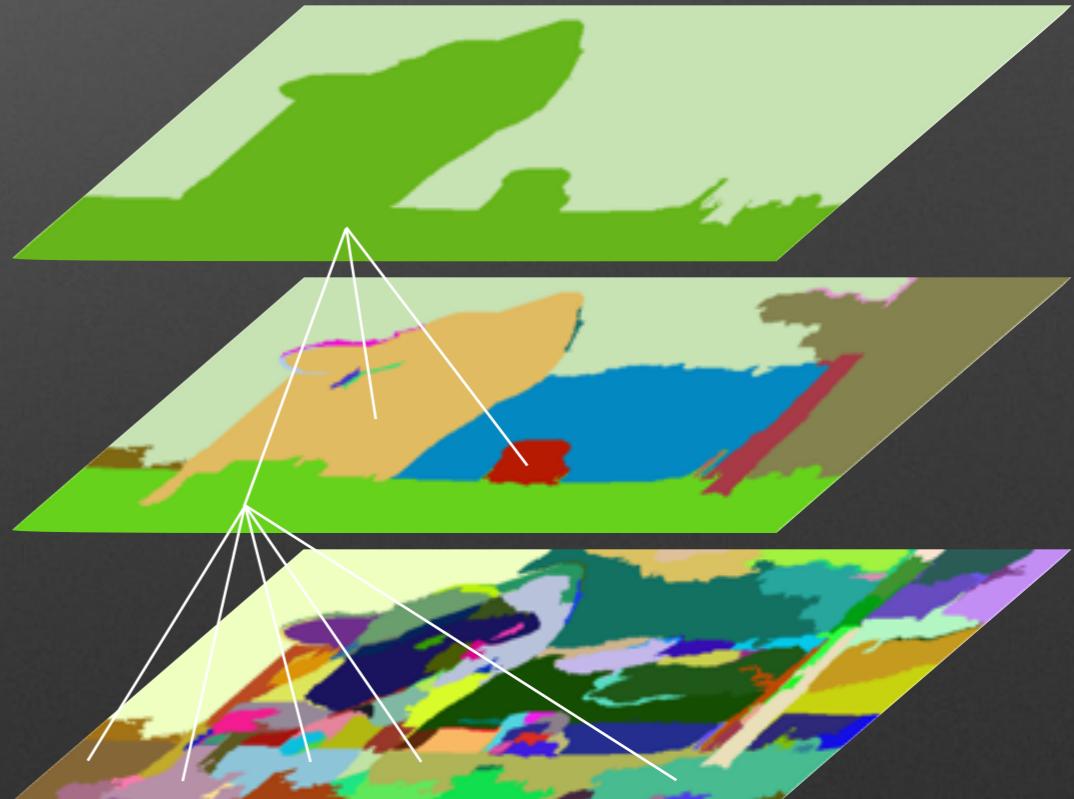
- Soup of segments
 - 3 different unsupervised segmentation algorithms
 - vary parameters
 - combination of up to 3 adjacent segments
- Slow (few minutes / image)
- Hard to segment larger objects well



Segmentation As Selective Search for Object Recognition

[van de Sande et al. 2011]

- Hierarchical segmentations
 - same algorithm
 - different parameters and color spaces
- Fast (few seconds / image)
- Good bounding box proposals
- OK segment proposals



Category Independent Object Proposals

[Endres and Hoiem 2010]

- Series of binary segmentations
 - Superpixel CRF model
 - Select seeds
 - Boosted decision tree as unary
 - Pairwise: boundary detector
 - Optimize with GraphCuts
- Good proposals
- Slow (100+ sec / image)



Constrained Parametric Min-Cuts for Automatic Object Segmentation

[Carreira and Sminchisescu 2010]

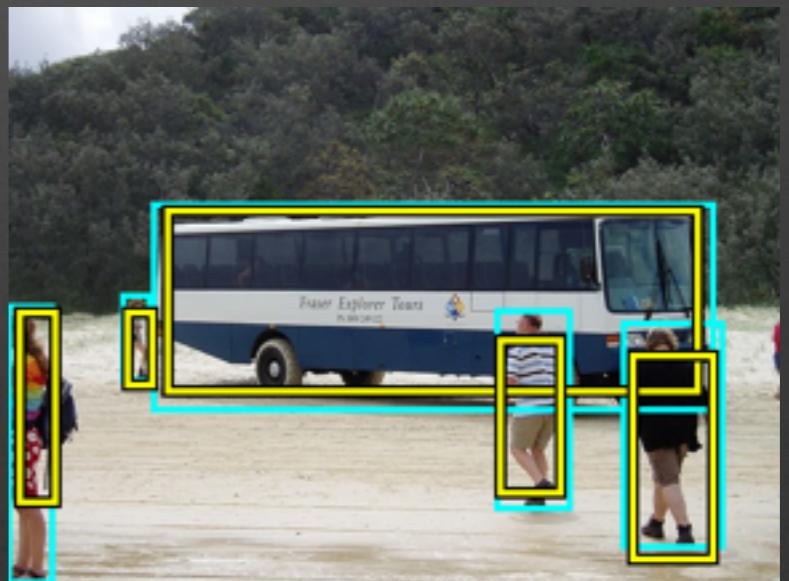
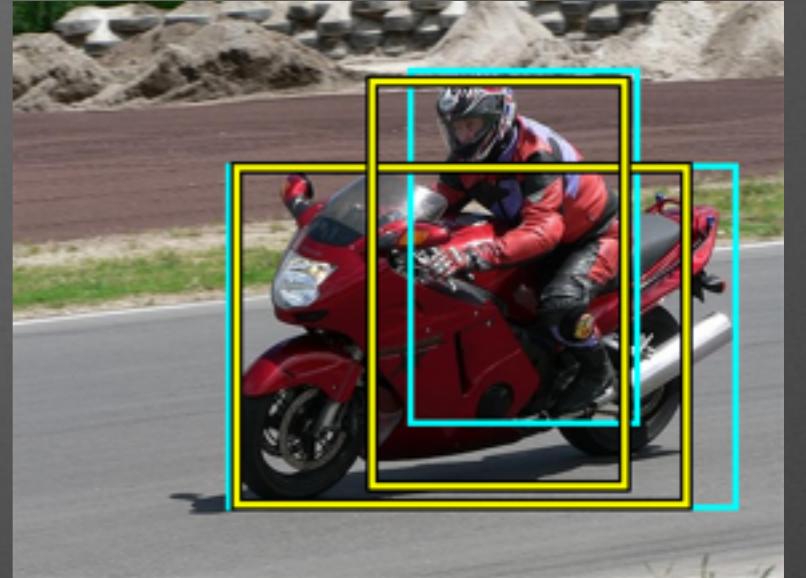
- Series of binary segmentations
 - Pixel CRF model
 - Regularly sampled seeds
 - Color based unary
 - GraphCuts with pairwise term
 - Parametric cut
 - slowly inflate segmentation
- Best proposals
- Slow (200+ sec / image)



What is an object?

[Alexe et al. 2010]

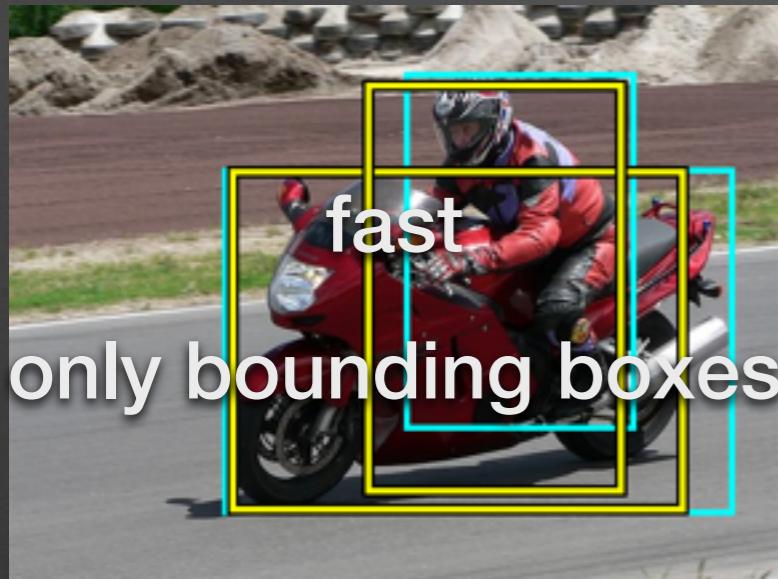
- Bounding box proposals
 - Train a general classifier for “objectness”
- fast (few seconds, BING¹: few ms)
- Only works for bounding boxes



[1] BING: Binarized Normed Gradients for Objectness Estimation at 300fps [Cheng et al. 2014]

Prior work - summary

Objectness



only bounding boxes

Segmentation based



good bounding boxes
OK segmentation

Seed / GraphCuts



Geodesic image segmentation

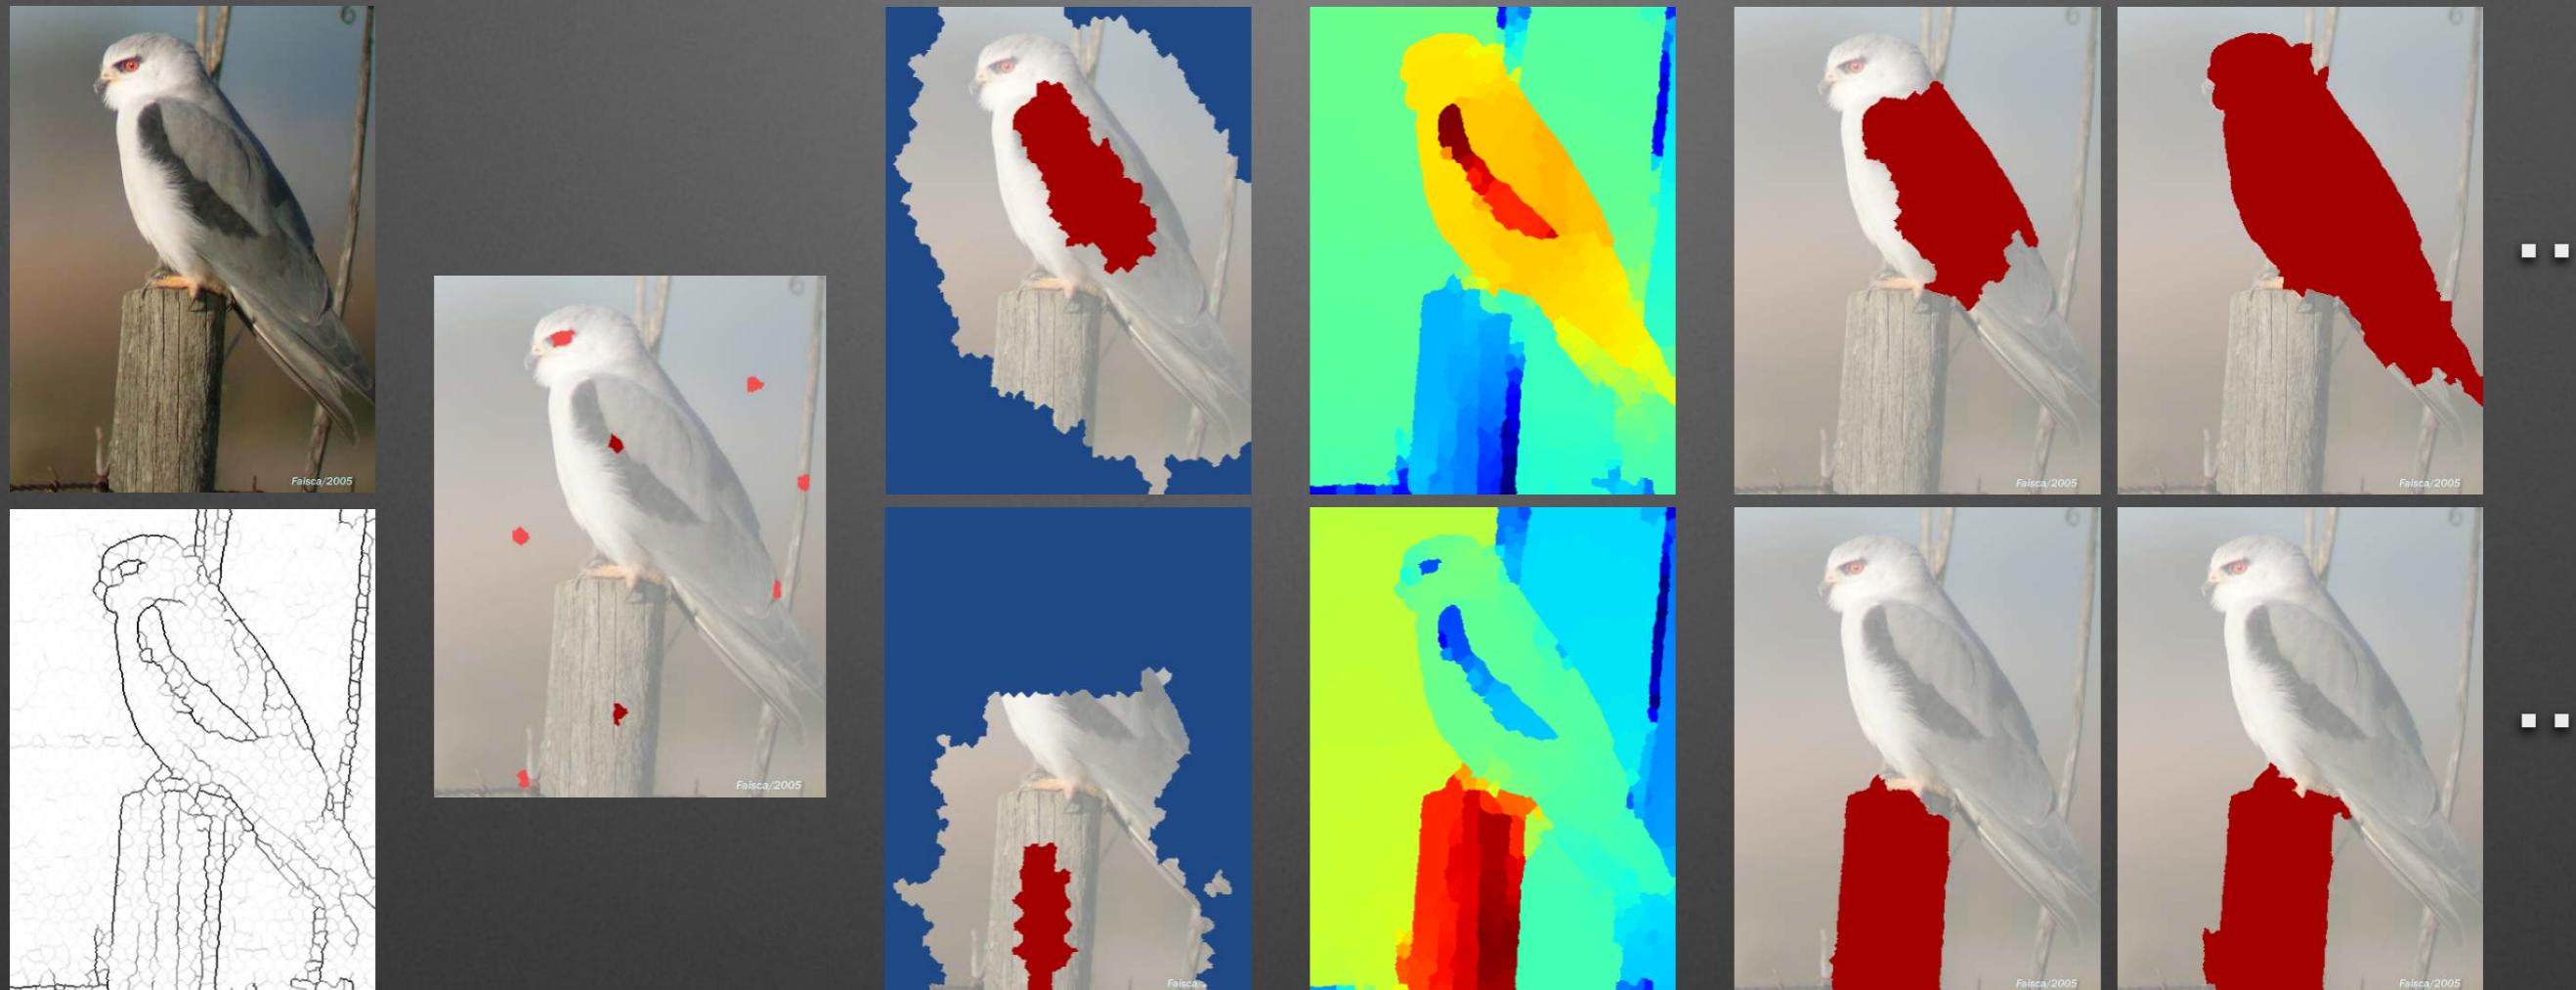
- (Signed) geodesic distance transform
 - Shortest path to **background** and **foreground** scribble
 - **small** within objects
 - **large** between objects
 - efficient to compute
- State of the art in interactive image and video segmentation^{1,2}



[1] Geodesic Matting: A Framework for Fast Interactive Image and Video Segmentation and Matting [Bai and Sapiro 2008]

[2] Geodesic Image and Video Editing [Criminisi et al. 2011]

Geodesic object proposals



image,
boundary
map and
superpixels

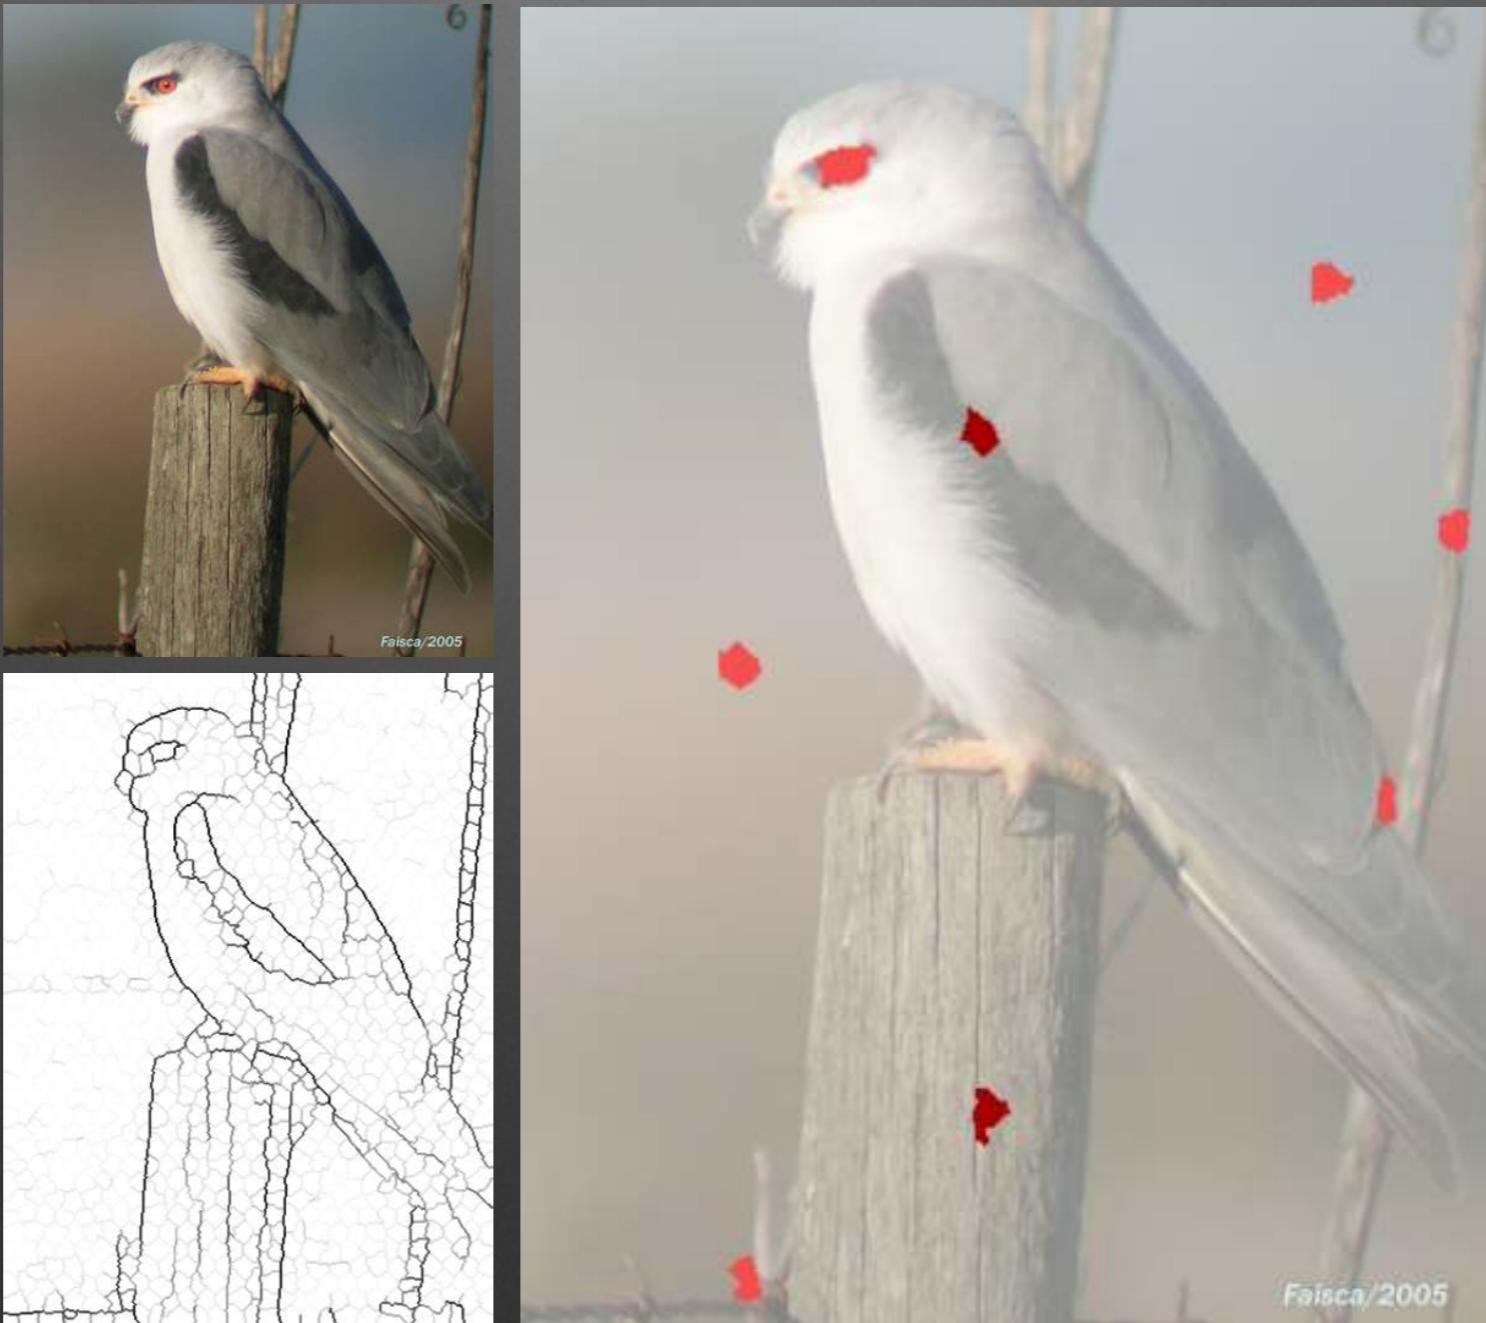
select
seeds

foreground
background
masks

geodesic
distance
transform

multiple proposals

Seed placement



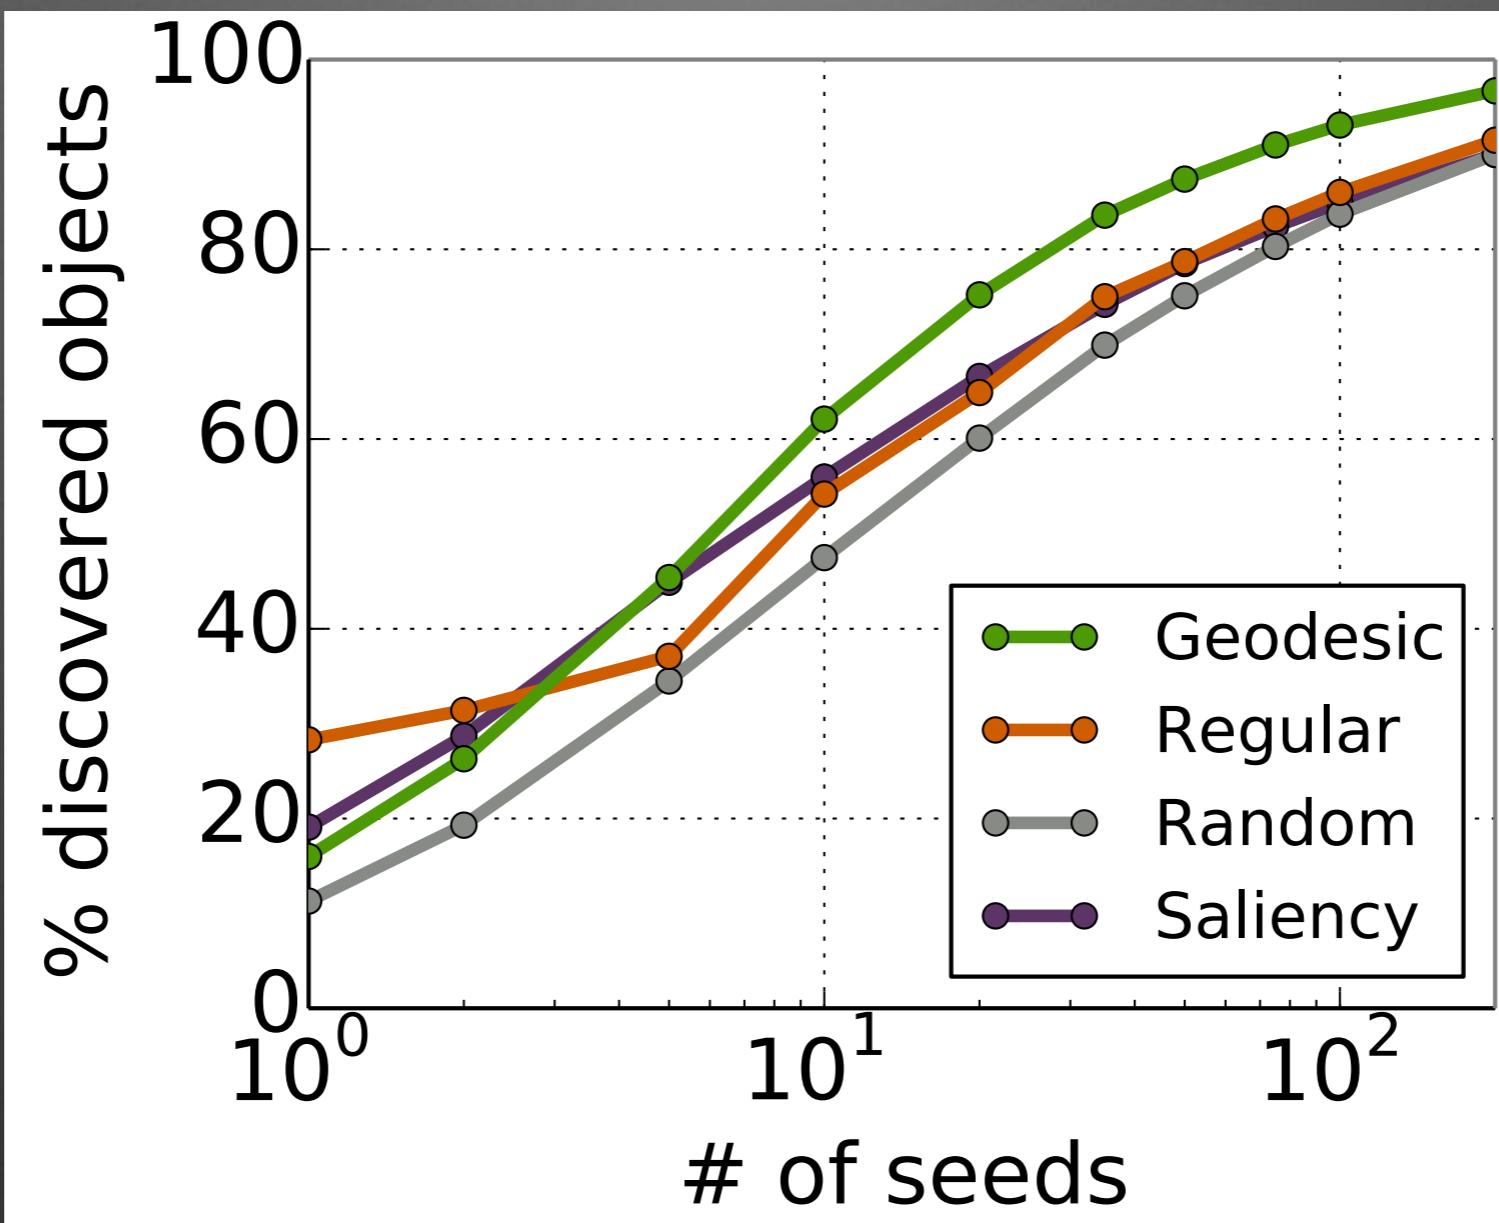
Seed placement

- Place a seed in each object
- Regular or random sampling
 - miss small objects
- Saliency based sampling
 - miss non-salient objects
- Geodesic placement
 - regular sampling in geodesic space
 - greedily place next seed at maximal geodesic distance



Faisca/2005

Seed placement

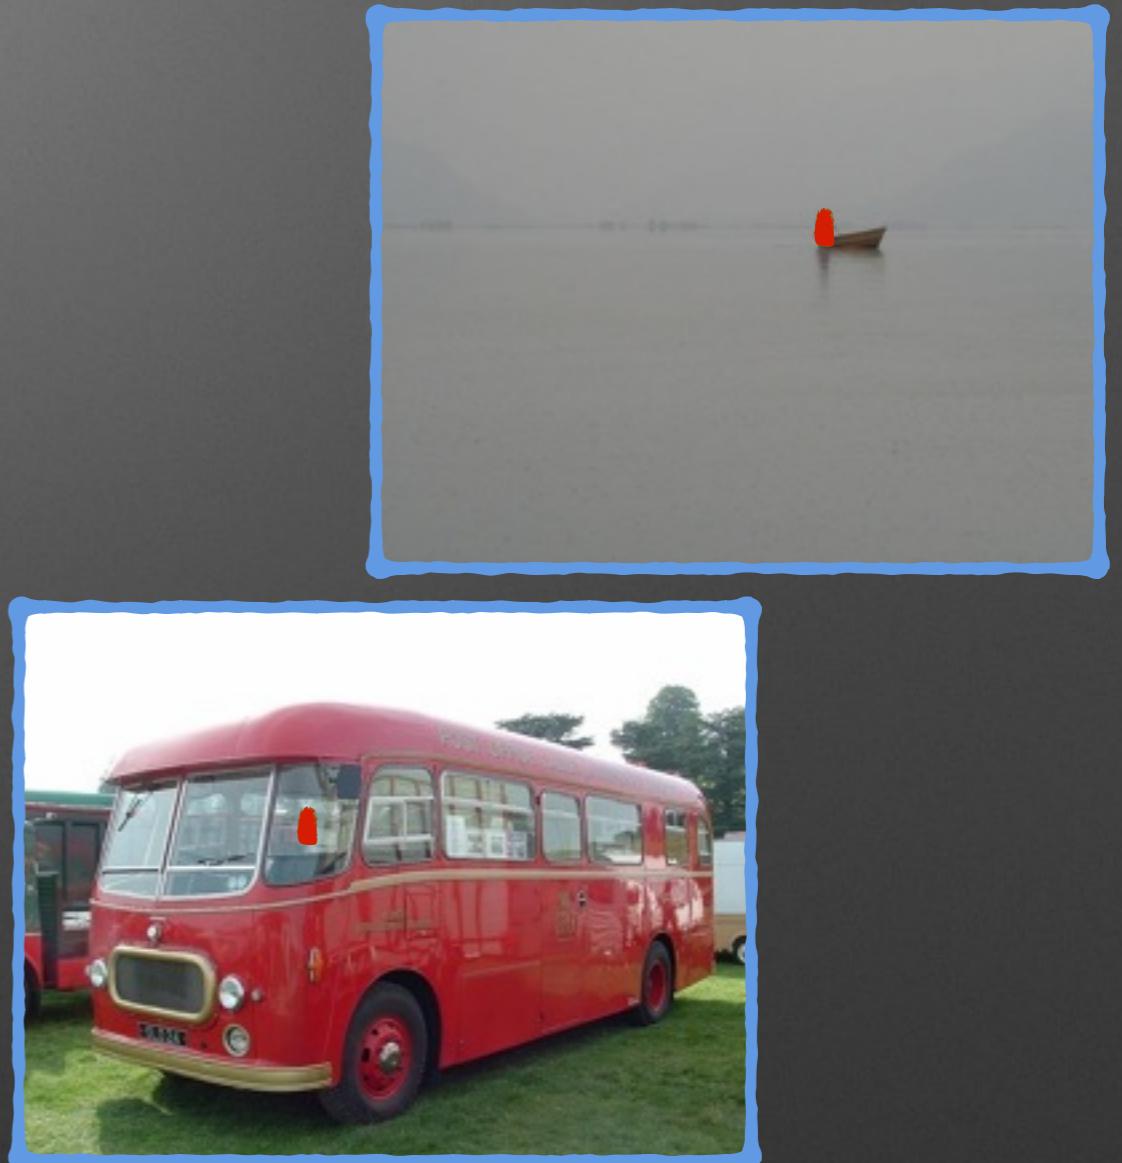


Mask generation



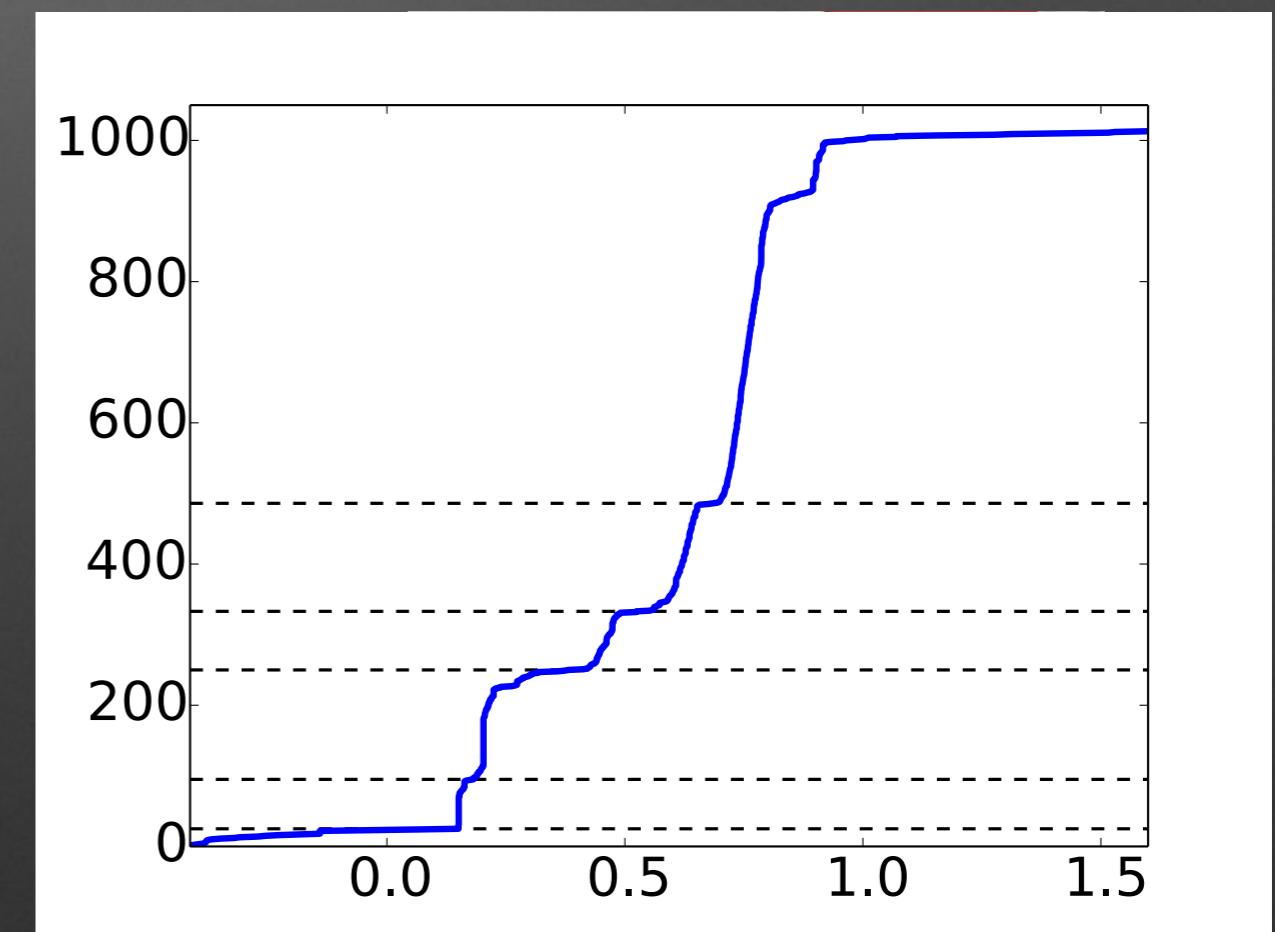
Mask generation

- No errors in masks
 - errors propagate
- Foreground mask
 - seed
- Background mask
 - boundary
 - empty

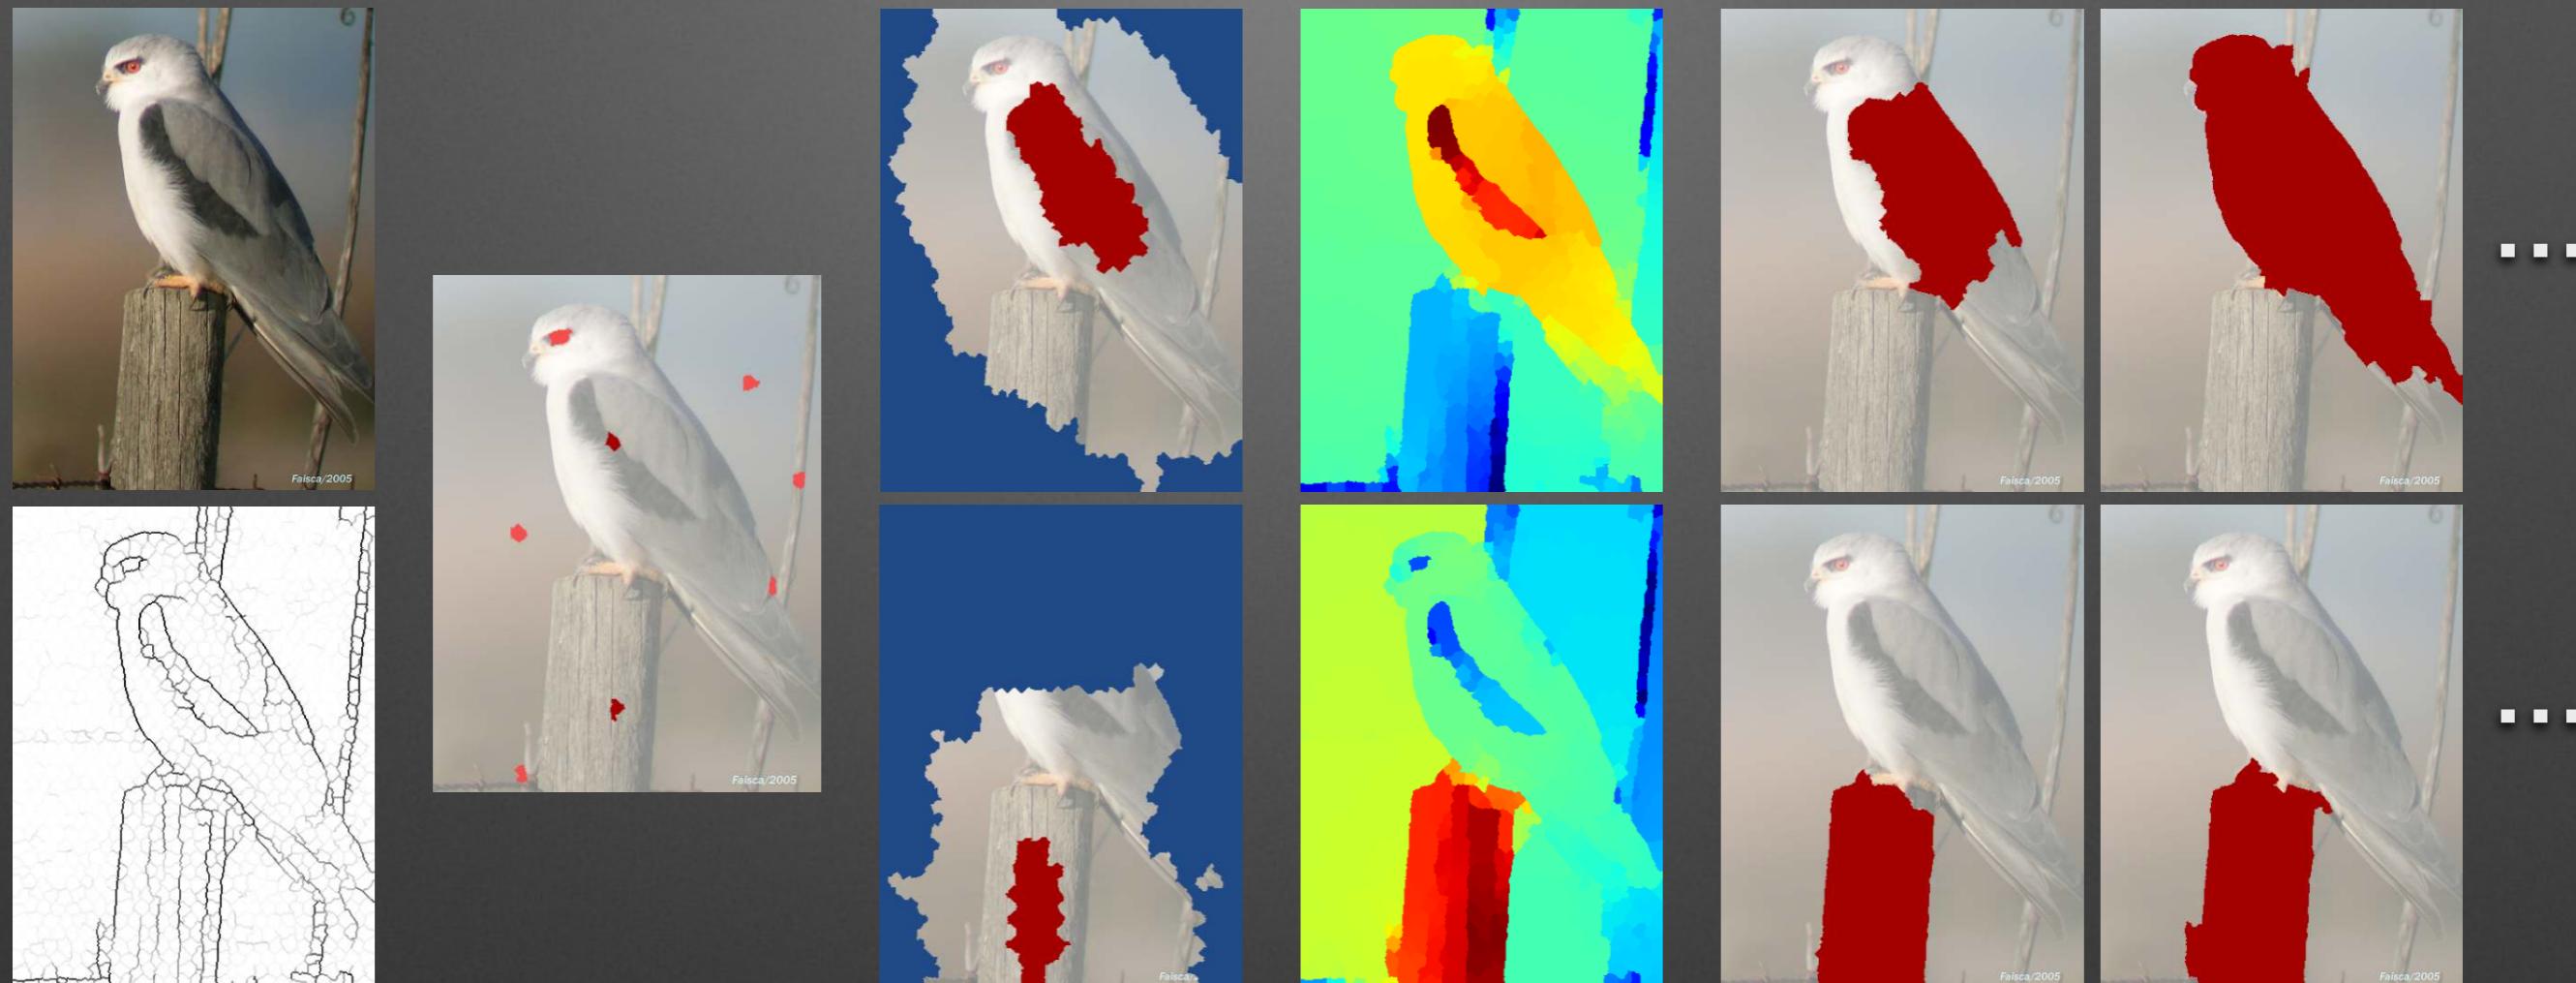


Geodesic segmentation

- Signed geodesic distance transform
- Each level set is a segmentation
- Find critical level sets
 - stationary points in geodesic function
 - evolution of Eikonal equation



Baseline GOP



image,
boundary
map and
superpixels

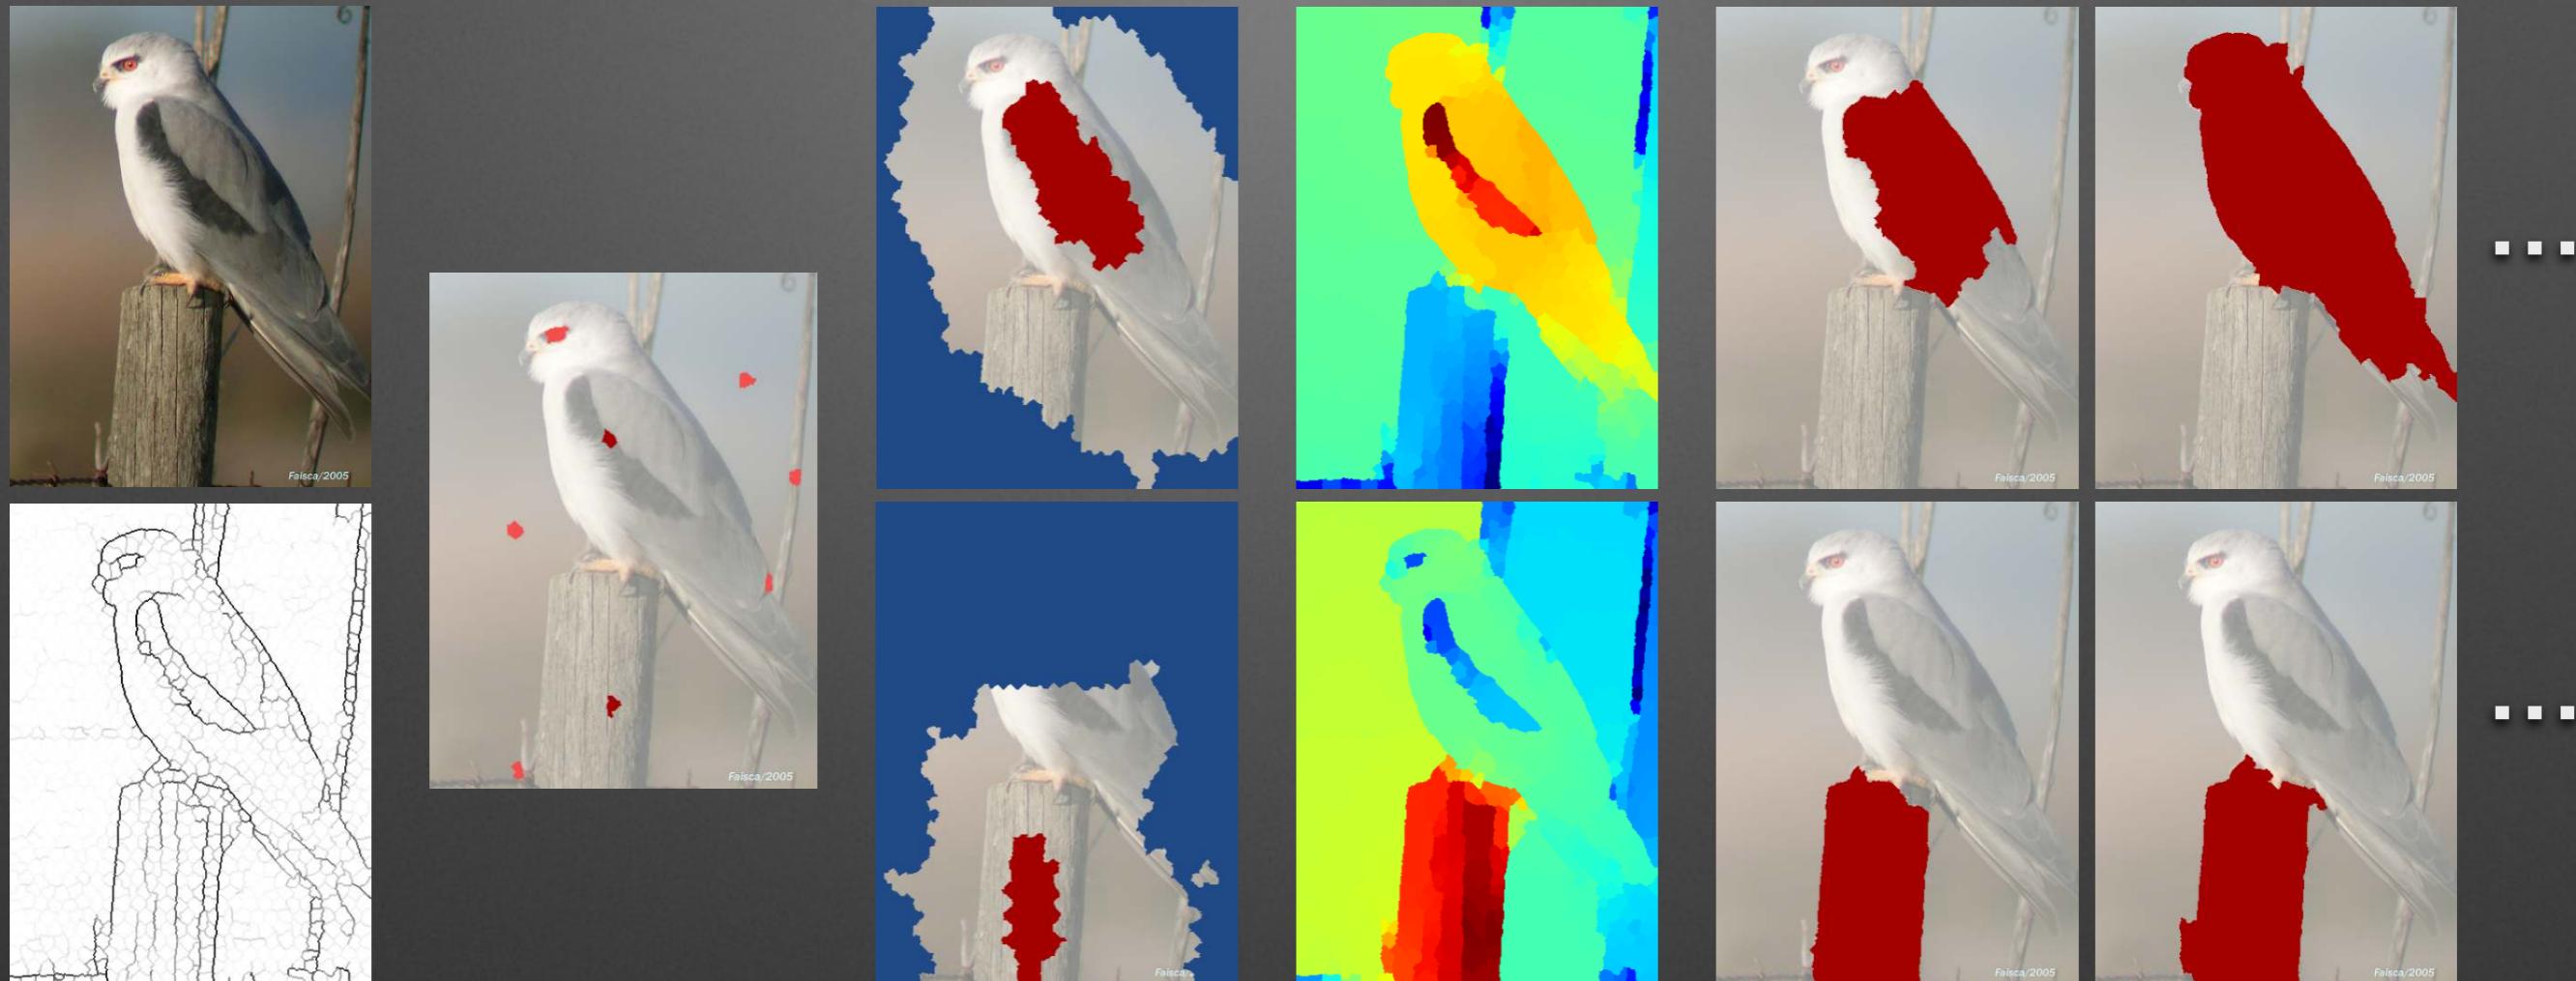
select
seeds

foreground
background
masks

geodesic
distance
transform

multiple proposals
per transform

Learned GOP



image,
boundary
map and
superpixels

learned
seeds

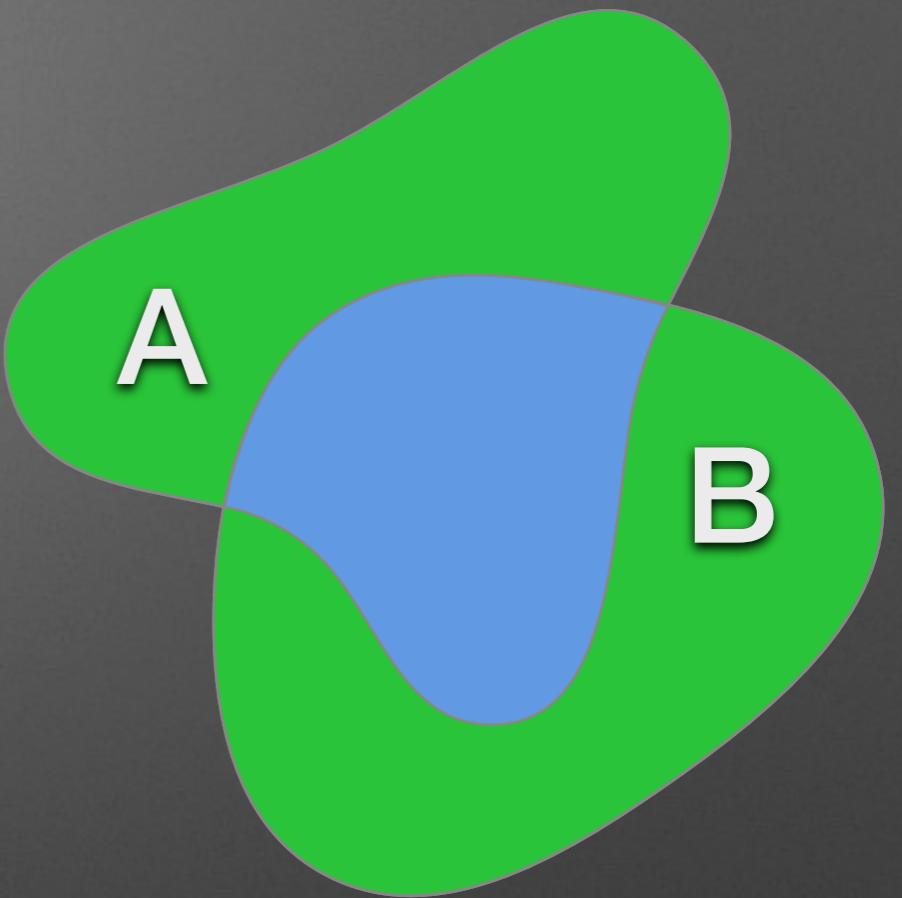
learned
masks

geodesic
distance
transform

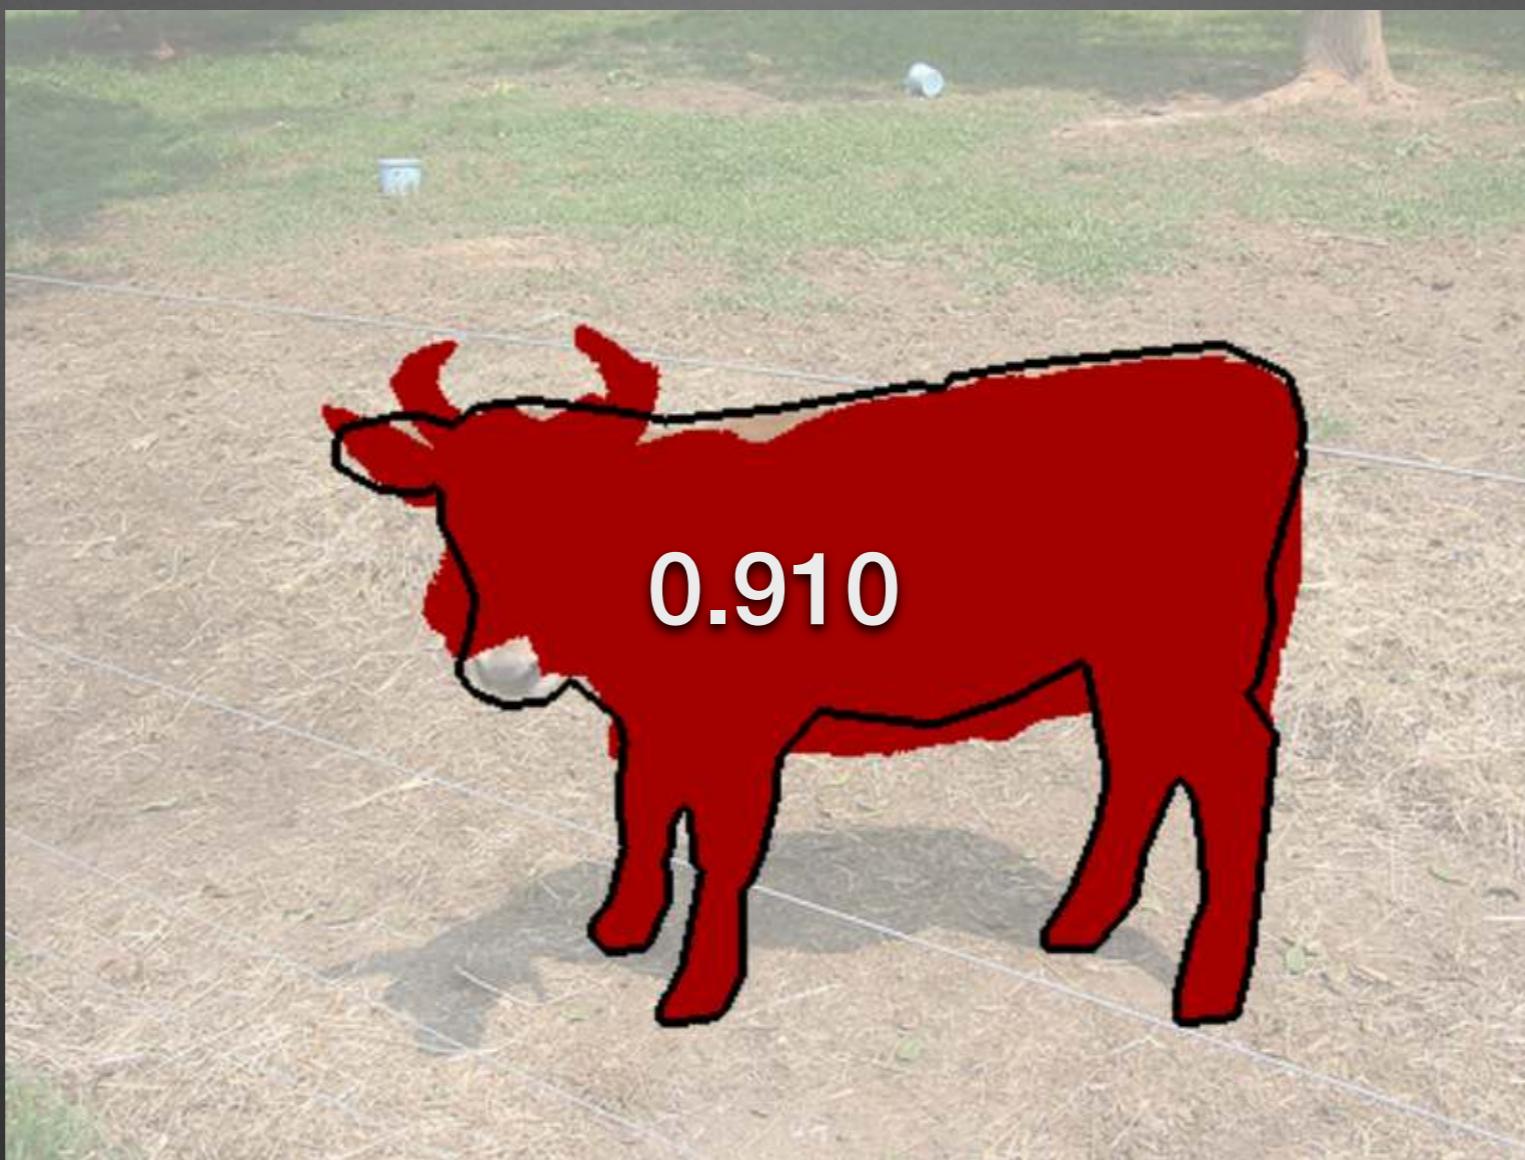
multiple proposals
per transform

Results

- VOC 2012 dataset
- Evaluation metric
 - **overlap** $\mathcal{J}(A, B) = \frac{|A \cap B|}{|A \cup B|}$
 - **best overlap** $b(O_k) = \max_P \mathcal{J}(O_k, P)$
 - **Average best overlap (ABO)**
$$\frac{1}{N} \sum_k b(O_k)$$
 - **α -recall**
$$\frac{1}{N} \sum_k [b(O_k) > \alpha]$$



What does overlap mean?

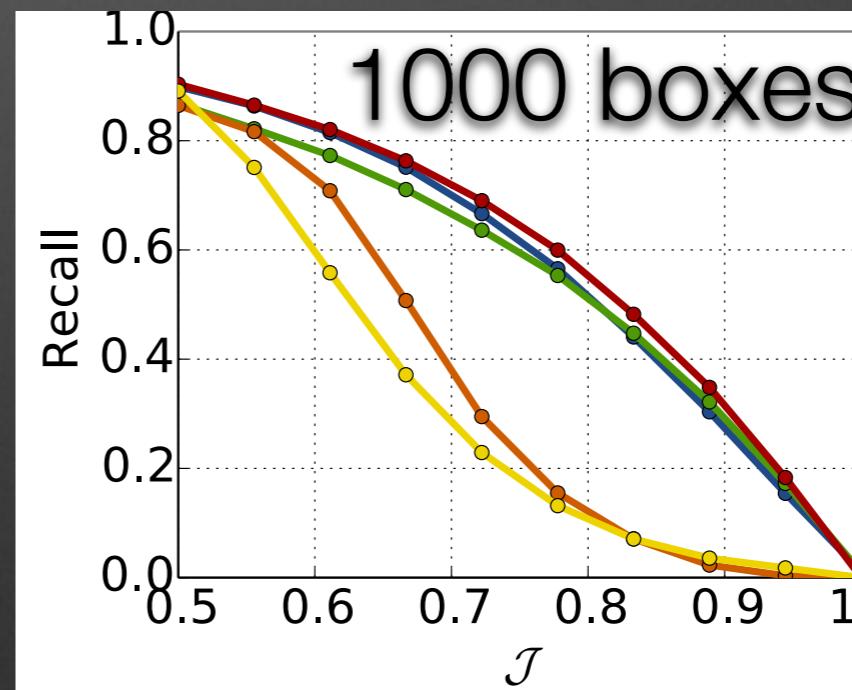
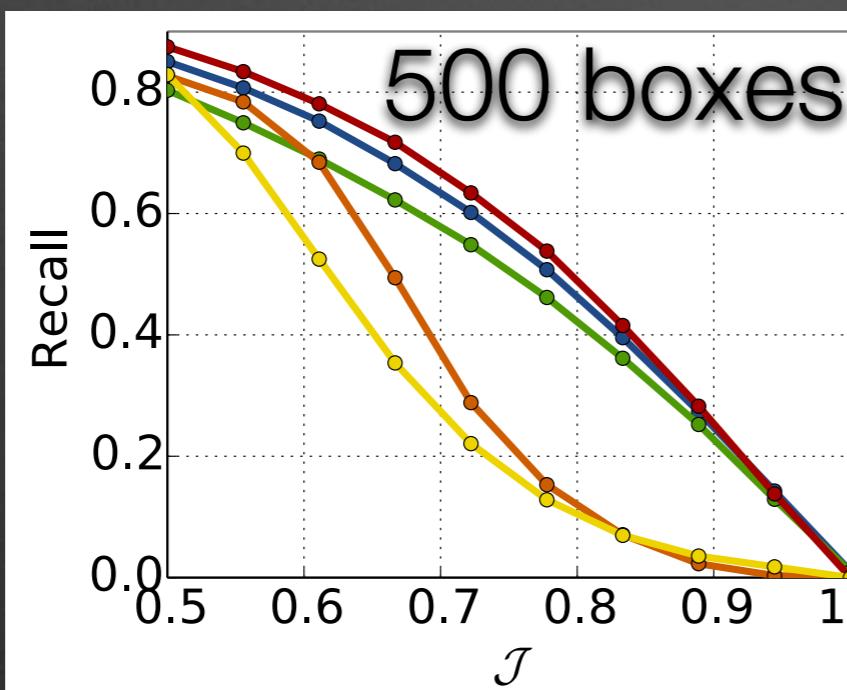
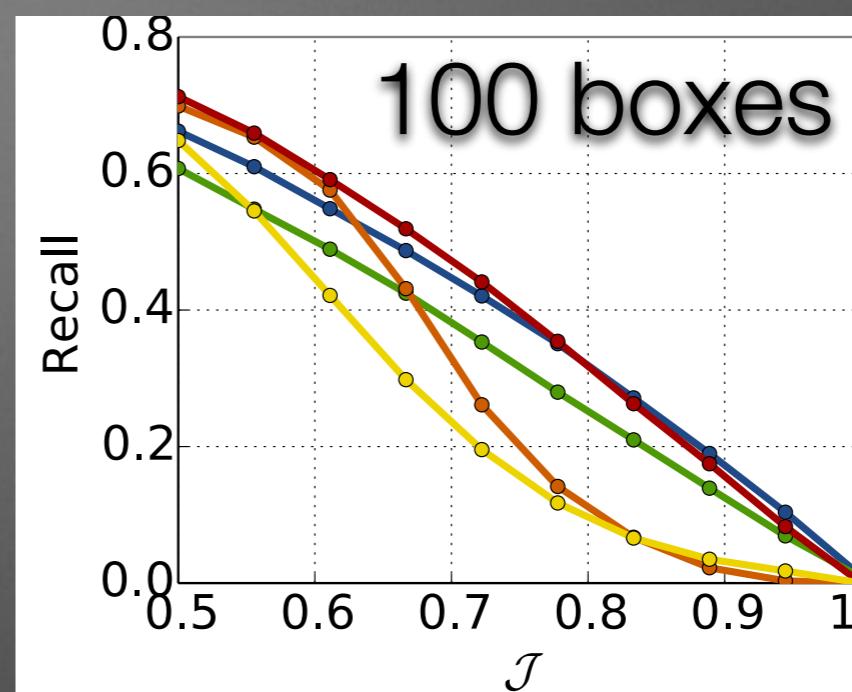
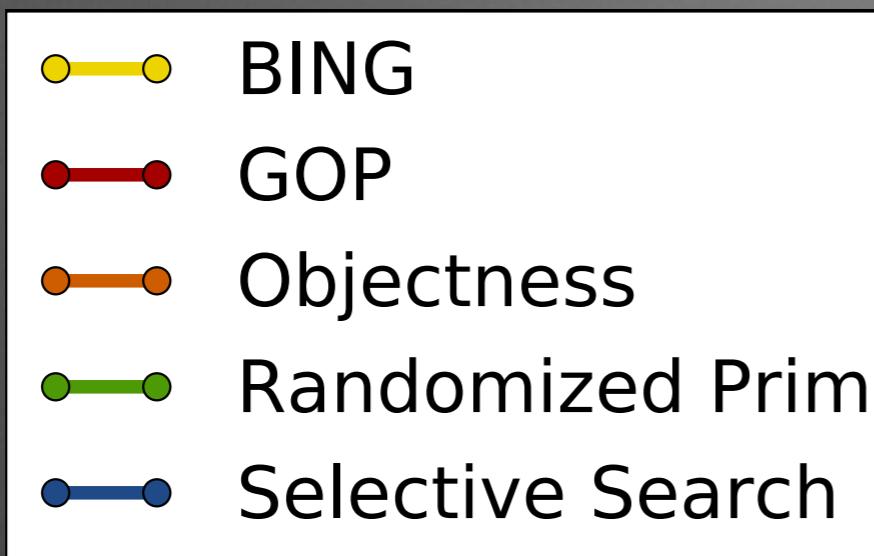


Segmentation results



METHOD	# PROP.	ABO	50%-RECALL	70%-RECALL	TIME
CPMC	646	0.703	0.784	0.609	252s
Baseline GOP	653	0.712	0.833	0.622	0.6s
Learned GOP	652	0.720	0.844	0.632	1.0s
Cat-Ind OP	1536	0.718	0.820	0.624	119s
Baseline GOP	1090	0.727	0.847	0.644	0.65
Learned GOP	1199	0.741	0.865	0.673	1.1s
Sel Search	4374	0.735	0.891	0.597	2.6s
Baseline GOP	2089	0.744	0.867	0.673	0.9s
Learned GOP	2286	0.756	0.877	0.699	1.4s
Baseline GOP	3958	0.756	0.881	0.699	1.2s
Learned GOP	4186	0.766	0.889	0.715	1.7s

Bounding box results



Bounding box results

	VOLUME UNDER SURFACE (VUS)
BING	0.278
Objectness	0.324
Randomized Prim	0.511
Selective Search	0.528
GOP	0.546

COCO dataset - segments



METHOD	# PROP.	50%-REC.	70%-REC.
Sel Search	6504	0.717	0.369
Baseline GOP	6106	0.704	0.426
Learned GOP	6264	0.717	0.447

LARGE OBJECTS ≥ 25			
METHOD	# PROP.	50%-REC.	70%-REC.
Sel Search	6504	0.810	0.442
Baseline GOP	6106	0.882	0.582
Learned GOP	6264	0.891	0.609

SMALL OBJECTS < 25			
METHOD	# PROP.	50%-REC.	70%-REC.
Sel Search	6504	0.525	0.219
Baseline GOP	6106	0.337	0.106
Learned GOP	6264	0.356	0.112



Summary

- Geodesic Object Proposals
 - fast
 - good segment proposals
 - good bounding box proposals
- Future work
 - small objects
 - learn proposals directly from data

Questions



C++, Python and Matlab Code:
<http://www.philkr.net/home/gop>

